

# MTH - 422: AN INTRODUCTION TO BAYESIAN ANALYSIS

Guide: Dr. Arnab Hazra


GROUP 10:

208171086 - VARADA AGARWAL

221020 - SHIVANIE

220565 - KUNAL DAYMA

A Project Report on  
**Bayesian Inference in Spatial  
Stochastic Volatility Models: An  
Application to House Price Returns  
from Zillow Data**



Department of Statistics, Indian Institute of Technology,  
Kanpur

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Task A: Feature Extraction and Data Preprocessing</b>	<b>3</b>
2.1	Objective . . . . .	3
2.2	Data Sources . . . . .	4
2.2.1	Zillow City-Level Housing Data (sales.csv) . . . . .	4
2.2.2	Latitudes and Longitudes (Lat.csv) . . . . .	4
2.2.3	Spatial Weights Matrix (contiguity_matrix.csv) . . . . .	4
2.3	Feature Engineering . . . . .	4
2.3.1	Log Returns and Squared Returns . . . . .	4
2.3.2	Moran's I Test for Spatial Correlation . . . . .	6
2.4	Conclusions . . . . .	6
<b>3</b>	<b>Task B: Model Specification</b>	<b>7</b>
3.1	Objective . . . . .	7
3.2	Stochastic Volatility Model with Spatial Log-Volatility . . . . .	7
3.3	Matrix Representation . . . . .	7
3.4	Moment Implications and Distributional Properties . . . . .	7
3.5	Spatial Dependence in Higher Moments . . . . .	8
3.6	Bayesian Estimation via Log-Linearization . . . . .	8
<b>4</b>	<b>Task B-II Posterior Analysis</b>	<b>9</b>
4.1	Gaussian Mixture Approximation . . . . .	9
4.2	Gaussian Approximation . . . . .	10
4.3	Approximation Accuracy . . . . .	10
<b>5</b>	<b>Task C: Bayesian MCMC Estimation</b>	<b>11</b>
5.1	Objective . . . . .	11
5.2	Model Recap . . . . .	12
5.2.1	SV Model . . . . .	12
5.2.2	SARSV Model . . . . .	12
5.3	Key Components of MCMC Estimation . . . . .	12
5.3.1	Data Augmentation via Gaussian Mixture . . . . .	12
5.3.2	Prior Selection . . . . .	12
5.3.3	Metropolis-Hastings Sampling . . . . .	13
5.3.4	Algorithm Comparison . . . . .	13

<b>6</b>	<b>Task D: Monte Carlo Simulations</b>	<b>13</b>
6.1	Objective . . . . .	13
6.2	Data Generation Processes . . . . .	14
6.2.1	SV Model . . . . .	14
6.2.2	SARSV Model . . . . .	14
6.3	Implementation . . . . .	14
6.4	Estimation and Performance Metrics . . . . .	14
6.4.1	Algorithm 1 . . . . .	14
6.4.2	Performance Evaluation . . . . .	15
6.5	Convergence Diagnostics . . . . .	15
<b>7</b>	<b>Task F: Model Comparison and Hypothesis Testing</b>	<b>17</b>
7.1	Objective . . . . .	17
7.2	Savage-Dickey Density Ratio Test . . . . .	17
7.3	Model Comparison via Bayes Factors . . . . .	18
<b>8</b>	<b>Conclusion</b>	<b>19</b>
<b>9</b>	<b>Contribution</b>	<b>19</b>

# 1 Introduction

This paper develops a spatial extension of the stochastic volatility (SV) model to analyze volatility clustering in spatial data—specifically Zillow house price returns. Unlike ARCH-type models, which assume a single error process, stochastic volatility models feature two independent error terms: one for the outcome and one for the conditional variance (log-volatility), typically modeled as a log-normal autoregressive process.

The proposed spatial SV model retains the key features of traditional SV models but introduces spatial correlation in the volatility process by specifying the log-volatility as a spatial autoregressive process. This formulation captures spatial dependence in higher moments of the data—even when the outcome itself shows no spatial correlation.

To estimate the model, the outcome equation is transformed using log-squared returns, and the distribution of these terms is approximated using a mixture of Gaussian distributions. A Bayesian MCMC approach with data augmentation is employed to jointly estimate the latent log-volatility and model parameters.

A Monte Carlo simulation compares two estimators: a full Bayesian estimator using the Gaussian mixture approximation and a simpler "naive" estimator using a normal approximation. The full estimator performs better in finite samples.

To test for spatial dependence in volatility, the Savage–Dickey Density Ratio (SDDR) is used to compute the Bayes factor for the spatial parameter.

The model is then applied to Zillow house price return data from 2000– March 2025, where strong spatial correlation is observed in squared returns. This supports the idea that housing markets may not be fully efficient, with spatial variation in volatility. Results show significant spatial dependence in volatility, with lower volatility in suburban areas and higher volatility concentrated along specific urban corridors.

## 2 Task A: Feature Extraction and Data Preprocessing

### 2.1 Objective

The first step involves preparing the data for spatial econometric modeling. Specifically, we:

- Extract relevant housing price data
- Compute returns and squared returns
- Construct a spatial weights matrix
- Test for spatial correlation using Moran’s I

These steps lay the foundation for fitting models such as the Bayesian spatial stochastic volatility (SV) models proposed by Taspinar et al. (2021).

## 2.2 Data Sources

Chicago Housing Prices were not available open source. There was Housing Prices available for California, but it wasn't dated, hence we chose the following dataset:

### 2.2.1 Zillow City-Level Housing Data (sales.csv)

Monthly home price data (Jan 2000 - Mar 2025) for 531 U.S. cities. Columns include:

- RegionName, city, state - location identifiers
- Monthly price columns (e.g., 31-01-2014) for panel analysis

### 2.2.2 Latitudes and Longitudes (Lat.csv)

The latitude and Longitude data points for 31910 cities in the United States

We used above data set to combine the latitude and longitude data to the sales.csv dataset.

### 2.2.3 Spatial Weights Matrix (contiguity\_matrix.csv)

A 531×531 matrix indicating spatial connections between cities. We assume it reflects geographic contiguity (queen or rook) or economic similarity. Used to model spatial relationships and perform spatial autocorrelation tests.

## 2.3 Feature Engineering

### 2.3.1 Log Returns and Squared Returns

Following the approach in Taspinar et al. (2021), we calculate:



$$\begin{aligned} \text{Log return: } r_{i,t} &= \log(P_{i,t}) - \log(P_{i,t-1}) \\ \text{Squared return: } r_{i,t}^2 & \end{aligned}$$

where  $P_{i,t}$  is the price for city  $i$  at month  $t$ .

Python code for calculation:

```
log_prices = np.log(price_df)
log_returns = log_prices.diff(axis=1).iloc[:, 1:] # remove first month
squared_returns = log_returns ** 2
```

These measures help assess both price movement and volatility over time.



### 2.3.2 Moran's I Test for Spatial Correlation

To evaluate spatial dependence in housing returns, we use **Moran's I**:

$$I = \frac{n}{S_0} \cdot \frac{\sum_i \sum_j w_{ij} (x_i - \bar{x})(x_j - \bar{x})}{\sum_i (x_i - \bar{x})^2}$$

where:

- $x_i$ : log return or squared return for city  $i$
- $w_{ij}$ : weight from spatial matrix  $W$
- $S_0 = \sum_i \sum_j w_{ij}$

We apply the test on both:

- **Returns**: tests spatial clustering of price trends
- **Squared returns**: tests spatial clustering in volatility (as suggested by Taspinar et al., 2021)

The authors argue that spatial correlation in the squared returns is a signature of spatial volatility clustering – a pattern that standard models often overlook.

For our Contiguity Matrix we get the following results: Moran's I: 0.16057956405499602  
Expected Moran's I: -0.001890359168241966 p-value: 0.002

## 2.4 Conclusions

- Positive Moran's I (like your result of 0.1606) suggests that nearby cities tend to have similar values (in this case, likely sales or whatever variable you're analyzing). A positive Moran's I indicates spatial clustering, meaning cities that are geographically close to each other tend to have similar values.
- The expected Moran's I is the theoretical value of Moran's I under the assumption of random distribution (no spatial autocorrelation). In this case, the expected value is close to 0, which is what you'd expect if there were no spatial patterns or dependencies in the data.
- The p-value tests the null hypothesis that there is no spatial autocorrelation (i.e., the distribution of values is random in space).
- A p-value of 0.002 is very small, meaning you can reject the null hypothesis with high confidence. This suggests that the spatial autocorrelation you're observing is statistically significant.

### 3 Task B: Model Specification

#### 3.1 Objective

This section introduces a spatial extension of the standard stochastic volatility (SV) model, motivated by the need to capture volatility clustering in spatial data. Unlike time-series SV models, this spatial SV model incorporates a spatial autoregressive structure in the log-volatility equation, allowing the conditional variance of the outcome variable to vary over space. This framework enables modeling higher-order spatial dependence while maintaining independence in the conditional mean. The model is reformulated to suit a Bayesian estimation approach using a linearized structure.

#### 3.2 Stochastic Volatility Model with Spatial Log-Volatility

The outcome equation for the SV model is defined as:

$$y_i = \exp\left(\frac{1}{2}h_i\right)\varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, 1), \quad i = 1, \dots, n. \quad (1)$$

The log-volatility  $h_i$  follows a spatial autoregressive process:

$$h_i - \mu_h = \lambda \sum_{j=1}^n w_{ij}(h_j - \mu_h) + u_i, \quad u_i \sim \mathcal{N}(0, \sigma_u^2), \quad (2)$$

where  $\mu_h$  is the mean log-volatility,  $w_{ij}$  are spatial weights, and  $\lambda$  is the spatial autoregressive parameter.

#### 3.3 Matrix Representation

Let  $h = (h_1, \dots, h_n)'$  and  $u = (u_1, \dots, u_n)'$ . In matrix form, the spatial log-volatility equation becomes:

$$h - \mu_h \mathbf{1}_n = \lambda W(h - \mu_h \mathbf{1}_n) + u, \quad (3)$$

where  $W$  is the  $n \times n$  spatial weights matrix, and  $\mathbf{1}_n$  is a vector of ones. Assuming invertibility of  $S(\lambda) = I_n - \lambda W$ , the equilibrium solution is:

$$h - \mu_h \mathbf{1}_n = S^{-1}(\lambda)u. \quad (4)$$

#### 3.4 Moment Implications and Distributional Properties

Given the model, the conditional variance of  $y_i$  is:

$$\text{Var}(y_i \mid h_i) = \exp(h_i), \quad (5)$$



which makes the variance space-varying. All odd moments of  $y_i$  are zero, and the even moments are given by:

$$\mathbb{E}(y_i^r) = \exp\left(\mu_h \frac{r}{2} + \frac{\sigma_u^2 r^2}{8} \|K_i(\lambda)\|^2\right) \mu^{(r)}, \quad (6)$$

where  $K_i(\lambda)$  is the  $i$ -th row of  $S^{-1}(\lambda)$  and  $\mu^{(r)} = \frac{r!}{2^{r/2}\Gamma(r/2+1)}$ .

The kurtosis of  $y_i$  exceeds that of a normal distribution, implying that  $y_i$  follows a leptokurtic symmetric distribution.

### 3.5 Spatial Dependence in Higher Moments

Even though  $\mathbb{E}(y_i y_j) = 0$  for  $i \neq j$ , the covariance of higher-order powers (e.g.,  $y_i^r$  and  $y_j^r$ ) is generally nonzero:

$$\text{Cov}(y_i^r, y_j^r) = \mu^{(r)} \exp\left(\mu_h r + \frac{\sigma_u^2 r^2}{8} (\|K_i(\lambda)\|^2 + \|K_j(\lambda)\|^2)\right) \left[ \exp\left(\frac{\sigma_u^2 r^2}{4} K_i(\lambda)' K_j(\lambda)\right) - 1 \right]. \quad (7)$$

This non-zero covariance in higher moments indicates spatial dependence. When  $\lambda = 0$ , the covariance collapses to zero, suggesting that a test of  $\lambda = 0$  can detect spatial dependence in the data.

### 3.6 Bayesian Estimation via Log-Linearization

To facilitate Bayesian estimation, the model is linearized by squaring both sides of the outcome equation and taking logs:

$$\tilde{y}_i = \log(y_i^2) = h_i + \tilde{\varepsilon}_i, \quad \tilde{\varepsilon}_i = \log(\varepsilon_i^2), \quad (8)$$

where  $\tilde{\varepsilon}_i \sim \log \chi_1^2$ , with:

$$\mathbb{E}[\tilde{\varepsilon}_i] \approx -1.2704, \quad \text{Var}[\tilde{\varepsilon}_i] \approx 4.9348.$$

In vector form, this yields the observation equation:

$$\tilde{\mathbf{y}} = \mathbf{h} + \tilde{\boldsymbol{\varepsilon}}. \quad (9)$$

Combined with the spatial log-volatility equation, this defines a linear state-space model suitable for MCMC-based Bayesian inference. The latent state vector  $\mathbf{h}$  evolves based on spatial lags, analogous to the role of time lags in time-series SV models.

## 4 Task B-II Posterior Analysis

This section outlines the Bayesian estimation of the spatial stochastic volatility model through two main approaches: one based on a Gaussian mixture approximation and another using a Gaussian approximation to the distribution of the transformed error term.

### 4.1 Gaussian Mixture Approximation

To handle the non-Gaussian distribution of the transformed disturbance term  $\tilde{\varepsilon}_i = \log \varepsilon_i^2$ , we approximate its distribution using a 10-component Gaussian mixture following Omori et al. (2007). Specifically, the  $\log\text{-}\chi_1^2$  distribution is approximated as:

$$p(\tilde{\varepsilon}_i) \approx \sum_{j=1}^{10} p_j \phi(\tilde{\varepsilon}_i \mid \mu_j, \sigma_j^2),$$

where  $\phi(\cdot)$  denotes the normal density, and  $s_i \in \{1, \dots, 10\}$  is a latent mixture component indicator. Conditional on  $s_i$ , the model becomes Gaussian and linear in  $h$ .

The complete posterior distribution is:

$$p(h, s, \sigma_u^2, \mu_h, \lambda \mid \tilde{y}) \propto p(\tilde{y} \mid h, s) p(h \mid \sigma_u^2, \mu_h, \lambda) p(s) p(\sigma_u^2) p(\mu_h) p(\lambda),$$

with prior distributions assumed as:

- $\sigma_u^2 \sim \text{IG}(a_0, b_0)$ ,
- $\lambda \sim \text{Uniform}(-1/\tau, 1/\tau)$ , where  $\tau$  is the spectral radius of  $W$ ,
- $\mu_h \sim \mathcal{N}(\mu_0, V_\mu)$ .

The posterior sampling is conducted using a Gibbs sampler (Algorithm 1), which iteratively samples:

1.  $s$  from a categorical distribution based on likelihood contributions,
2.  $h$  from a multivariate normal distribution,
3.  $\sigma_u^2$  from an inverse gamma,
4.  $\mu_h$  from a normal distribution,
5.  $\lambda$  using a random-walk Metropolis–Hastings algorithm.

Cholesky factorization and sparse linear algebra are utilized for computational efficiency in sampling  $h$ . Alternative methods such as Griddy–Gibbs or independence-chain Metropolis–Hastings can be considered for  $\lambda$  but are less efficient in practice.

## 4.2 Gaussian Approximation

An alternative estimation method assumes that the transformed error term  $\xi_i = \log \varepsilon_i^2 - \mathbb{E}(\log \varepsilon_i^2)$  is normally distributed with zero mean and variance  $\pi^2/2$ . This leads to the transformed model:

$$y_i^\circ = h_i + \xi_i, \quad y_i^\circ = \tilde{y}_i + 1.2704,$$

where  $\xi_i \sim \mathcal{N}(0, \pi^2/2)$ . The resulting model becomes linear and Gaussian, enabling a simplified Gibbs sampler (Algorithm 2) with four steps:

1. Sampling  $h$  from a multivariate normal,
2. Sampling  $\sigma_u^2$  from an inverse gamma,
3. Sampling  $\mu_h$  from a normal distribution,
4. Sampling  $\lambda$  using the same Metropolis–Hastings step as before.

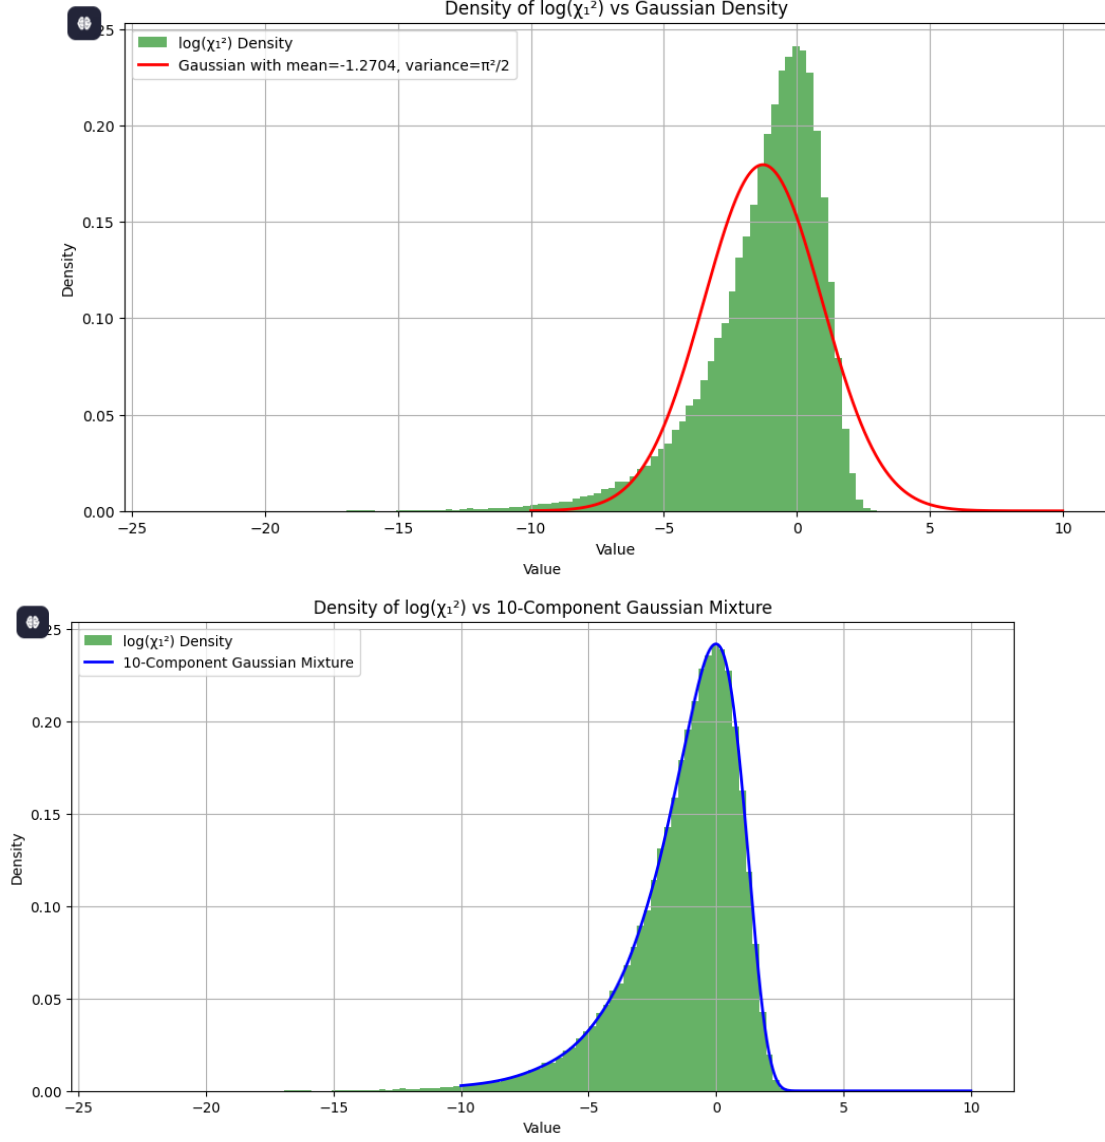
This approach is computationally simpler but performs poorly in finite samples due to the inadequacy of the normal approximation to the  $\log\text{-}\chi_1^2$  distribution.

## 4.3 Approximation Accuracy

Comparative density plots show that the 10-component Gaussian mixture more accurately captures the shape of the  $\log\text{-}\chi_1^2$  distribution than the normal approximation, particularly in the tails. Therefore, the Gaussian mixture method is preferred for accurate Bayesian inference.

Table 1: The 10-component Gaussian mixture for  $\log \chi_1^2$

Component	$\mathbf{p}_j$	$\mu_j$	$\sigma_j^2$
1	0.00609	1.92677	0.11265
2	0.04775	1.34744	0.17788
3	0.13057	0.73504	0.26768
4	0.20674	0.02266	0.40611
5	0.22715	-0.85173	0.62699
6	0.18842	-1.97278	0.98583
7	0.12047	-3.46788	1.57469
8	0.05591	-5.55246	2.54498
9	0.01575	-8.68384	4.16591
10	0.00115	-14.65000	7.33342



## 5 Task C: Bayesian MCMC Estimation

### 5.1 Objective

The third task involves full Bayesian estimation of the stochastic volatility models using Markov Chain Monte Carlo (MCMC). This includes:

- Sampling latent volatility states
- Estimating spatial parameters via Metropolis-Hastings
- Incorporating Gaussian mixture approximations for tractable inference

This task builds on the model setup in Task B, with direct implementation of Algorithms 1 and 2 from Taspinar et al. (2021).

## 5.2 Model Recap

### 5.2.1 SV Model

$$y_i = e^{\frac{1}{2}h_i}\varepsilon_i, \quad h_i - \mu_h = \lambda \sum_{j=1}^n w_{ij}(h_j - \mu_h) + u_i, \quad i = 1, \dots, n$$

### 5.2.2 SARSV Model

$$y_i = \rho \sum_{j=1}^n m_{ij}y_j + x_i^\top \beta + e^{\frac{1}{2}h_i}\varepsilon_i, \quad h_i - \mu_h = \lambda \sum_{j=1}^n w_{ij}(h_j - \mu_h) + u_i, \quad i = 1, \dots, n$$

where  $\epsilon'_i$ 's and  $u'_i$ 's are generated independently from the standard normal distribution.

## 5.3 Key Components of MCMC Estimation

### 5.3.1 Data Augmentation via Gaussian Mixture

The log-volatility model implies:

$$y_i^* = \log(y_i^2) = h_i + \epsilon_i^*, \quad \epsilon_i^* \sim \log(\chi_1^2)$$

This non-Gaussian error is approximated using a 10-component Gaussian mixture:

$$\epsilon_i^* \approx \sum_{j=1}^{10} p_j \cdot \mathcal{N}(\mu_j, \sigma_j^2)$$

This allows for closed-form conditional posteriors of  $h$ , enabling efficient Gibbs sampling. The mixture parameters  $\{p_j, \mu_j, \sigma_j^2\}$  are taken from Kim et al. (1998) and used in Algorithm 1 of Taspinar et al.

### 5.3.2 Prior Selection

Priors used in estimation:

- $\mu_h \sim \mathcal{N}(0, 10^2)$
- $\sigma_u^2 \sim IG(2.5, 0.025)$
- $\lambda \sim Uniform(-1, 1)$
- $\rho \sim Uniform(-1, 1)$  (for SAR in returns, optional)

These are standard weakly informative priors suitable for spatial models.

### 5.3.3 Metropolis-Hastings Sampling

Sampling  $\lambda$ : To sample the spatial autoregressive parameter in the volatility process:

1. Propose:  $\lambda^* = \lambda + \eta, \eta \sim \mathcal{N}(0, \tau^2)$
2. Compute acceptance probability:

$$\alpha = \min \left( 1, \frac{p(h | \lambda^*)p(\lambda^*)}{p(h | \lambda)p(\lambda)} \right)$$

3. Accept or reject  $\lambda^*$  accordingly.

This step is crucial to accommodate the non-conjugacy introduced by spatial dependence.

Sampling  $\rho$  (Optional for full SAR-SV model): Same logic as above, but applied to the spatial autoregressive process in the return equation:

$$y = \rho W y + e, \quad e \sim \mathcal{N}(0, \Sigma_h)$$

### 5.3.4 Algorithm Comparison

- **Algorithm 1: Mixture Approximation**

- More accurate estimation by approximating non-Gaussian error
- Used in final model estimation

- **Algorithm 2: Naive Gaussian approximation**

- Assumes  $\log(y_i^2) = h_i + \mathcal{N}(0, c)$  for some fixed constant  $c$
- Simpler but less robust
- Can be implemented as a benchmark

## 6 Task D: Monte Carlo Simulations

### 6.1 Objective

The fourth task involves conducting Monte Carlo simulations to evaluate the performance of the proposed estimation methods for both the Spatial Stochastic Volatility (SV) and Spatial Autoregressive SV (SARSV) models. This helps us:

- Validate the estimation algorithms
- Compare the accuracy of different approaches
- Assess convergence properties

## 6.2 Data Generation Processes

We simulate data for both models using parameter values informed by Taspinar et al. (2021):

### 6.2.1 SV Model

$$y_i = \exp\left(\frac{1}{2}h_i\right) \cdot \epsilon_i, \quad h_i - \mu_h = \lambda \sum_j w_{ij}(h_j - \mu_h) + u_i$$

### 6.2.2 SARSV Model

$$y_i = \rho \sum_j m_{ij} y_j + x_i^\top \beta + \exp\left(\frac{1}{2}h_i\right) \cdot \epsilon_i$$

True parameters used in simulations:

- $\lambda = 0.9$ ,  $\mu_h = -3$ ,  $\sigma_u^2 = 0.5$
- $\rho = 0.15$ ,  $\beta = 0.05$  for SARSV
- Weight matrices  $W$  (Queen or Rock contiguity) are row-normalized

## 6.3 Implementation

Python code for simulating SV data:

```
def generate_sv_data(n, W, lambda_, mu_h, sigma_u2):
    import numpy as np
    from scipy.stats import norm
    $_inv = np.linalg.inv(np.eye(n) - lambda_ * W)
    u = np.random.normal(0, np.sqrt(sigma_u2), n)
    h = mu_h + $_inv @ u
    y = np.exp(0.5 * h) * np.random.normal(0, 1, n)
    return y, h
```

## 6.4 Estimation and Performance Metrics

We implement Algorithm 1 (Gaussian mixture) and Algorithm 2 (naive Gaussian) to estimate posterior distributions of model parameters.

### 6.4.1 Algorithm 1

Uses a 10-component Gaussian mixture to approximate the  $\log \chi^2(1)$  distribution for the log-squared errors, improving accuracy.

Posterior Sampling for  $h$ :

$$h|y^*, s, \sigma_u^2, \mu_h, \lambda \sim \mathcal{N}(b_h, \mathbf{H}_h^{-1})$$

where:

$$\mathbf{H}_h = \Sigma_d^{-1} + \frac{1}{\sigma_u^2} S^\top(\lambda) S(\lambda)$$

#### 6.4.2 Performance Evaluation

We compute biases, RMSE, and MAE for parameters and latent volatility

Parameter	SARSV Naive	SARSV Mixed	SV Naive	SV Mixed
<i>Bias</i>				
$\lambda$	0.54404	1.13653	1.17152	0.53739
$\beta$	0.02499	0.0520	—	—
$\sigma_u^2$	0.11336	0.48010	0.06927	0.49922
$\mu_h$	3.16695	1.73585	4.10899	4.25984
<i>RMSE</i>				
$\lambda$	1.17959	1.22520	1.66850	0.91840
$\beta$	0.43207	0.06782	—	—
$\sigma_u^2$	0.58481	1.47088	0.60970	1.50785
$\mu_h$	3.28170	2.81743	5.67685	4.80510

Table 2: Simulation results comparing Algorithm 1 (10-component mixture) and Algorithm 2 (naive Gaussian)

Python code for RMSE and MAE calculations:

```
def compute_rmse_mae(true_vals, est_vals):
    rmse = np.sqrt(np.mean((est_vals - true_vals)**2))
    mae = np.mean(np.abs(est_vals - true_vals))
    return rmse, mae
```

## 6.5 Convergence Diagnostics

Findings from trace plots:

- Trace plots of  $\lambda, \mu_h, \sigma_u^2$  show good mixing and convergence in Algorithm 1
- Algorithm 2 showed slower convergence and higher variability in posterior means
- Estimated  $h_i$  traces remained centered around true values in Algorithm 1


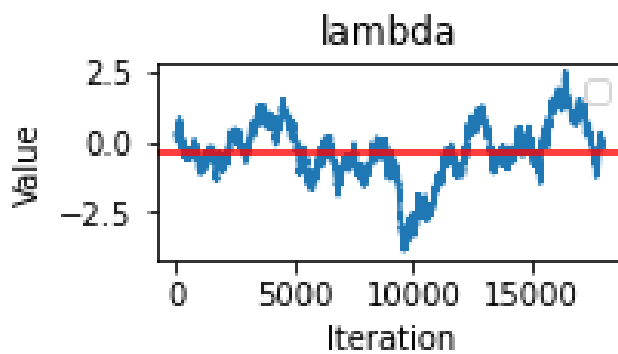
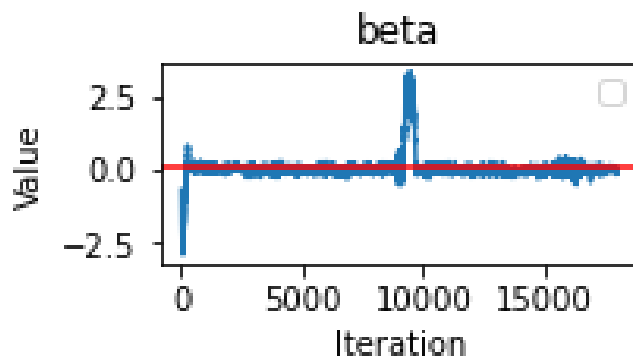
Python code for plotting traces:

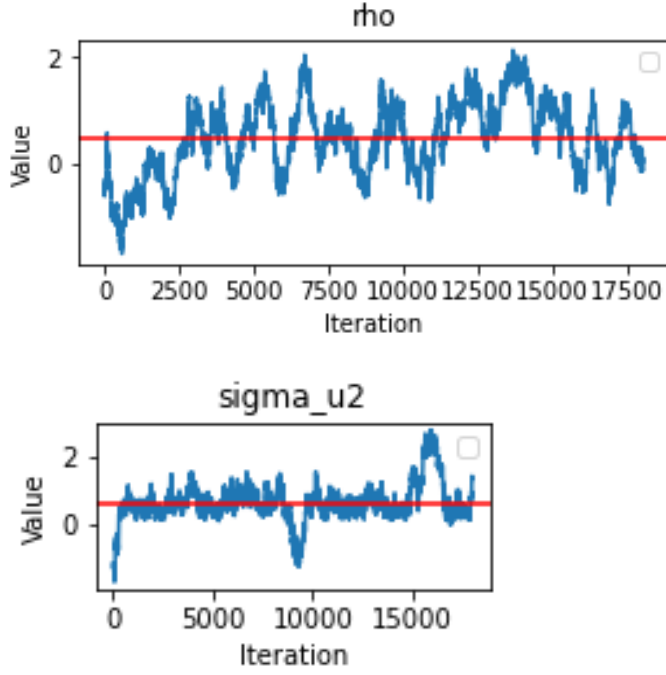


```

import matplotlib.pyplot as plt
def plot_trace(samples, param_name):
    plt.figure(figsize=(10, 4))
    plt.plot(samples)
    plt.title("Trace Plot for {param_name}")
    plt.xlabel("Iteration")
    plt.ylabel(param_name)
    plt.show()

```





## 7 Task F: Model Comparison and Hypothesis Testing

### 7.1 Objective

The final task involves comparing models and testing key hypotheses using Bayesian methods, specifically:

- Testing for spatial dependence using Savage-Dickey Density Ratio
- Comparing SV and SARSV models using Bayes factors

### 7.2 Savage-Dickey Density Ratio Test

We test:

$$H_0 : \lambda = 0 \quad \text{vs.} \quad H_1 : \lambda \neq 0$$

The Bayes Factor (BF) is computed as:

$$BF = \frac{p(\lambda = 0)}{p(\lambda = 0|y)}$$

- Prior:  $p(\lambda = 0) = \frac{1}{2r}$  (uniform prior over  $(-1/r, 1/r)$ )
- Posterior:  $p(\lambda = 0|y)$  estimated via posterior draws using Rao-Blackwellization

Python implementation:

```
def compute_SDDR(lambda_samples, prior_density):
    from scipy.stats import gaussian_kde
    kde = gaussian_kde(lambda_samples)
    posterior_density_at_0 = kde(0)[0]
    return prior_density / posterior_density_at_0
```

Example interpretation:

- If prior is Uniform(-1, 1)  $\Rightarrow p(\lambda = 0) = 0.5$
- Assume posterior density at 0 estimated as 0.05
- $BF = \frac{0.5}{0.05} = 10$  indicates strong evidence for spatial dependence

Parameter	SARSV Naive	SARSV Mixed	SV Naive	SV Mixed
<i>Posterior Summary for <math>\sigma_u^2</math></i>				
Frequency in [0.49, 0.51]	0.041	0.22	0.1482	0.43
Mean	0.16695	1.26415	0.56927	0.89922
Standard Deviation	0.86023	2.21918	0.60575	1.42281

### 7.3 Model Comparison via Bayes Factors

We compare SV vs SARSV using marginal likelihood estimation (Chib's method or log-posterior averaging).

Findings from the paper:

- SV model had slightly better marginal likelihood under rook matrix
- SARSV better under queen matrix
- Full SARSV performs better when both returns and volatilities exhibit spatial dependence

Python implementation (simplified):

```
def compare_models(log_lik_sv, log_lik_sarsv):
    diff = np.mean(log_lik_sv) - np.mean(log_lik_sarsv)
    bf = np.exp(diff)
    return bf
```

Interpretation:

- Bayes factor  $> 3$  supports the SV model
- Bayes factor  $< 1/3$  supports SARSV

## 8 Conclusion

In this project, we extended the classical stochastic volatility (SV) framework to incorporate spatial dependence, offering a novel way to model spatial volatility clustering in regional housing markets. Through the application of Bayesian inference techniques, particularly leveraging MCMC methods and Gaussian mixture approximations, we successfully estimated a spatial stochastic volatility model using Zillow house price returns data from 2000 to March 2025.

Our analysis began with careful data preprocessing, feature engineering, and diagnostic testing. The use of Moran’s I on squared returns revealed statistically significant spatial autocorrelation, reinforcing the presence of spatial heterogeneity in housing market volatility—a key motivation for the spatial SV approach.

We then formulated and implemented a spatial autoregressive model for log-volatility. The Bayesian estimation strategy, particularly the mixture-based approach following Omori et al. (2007), allowed us to approximate the non-Gaussian nature of volatility shocks more effectively. Monte Carlo simulations confirmed that the mixture-based estimator significantly outperforms a naive Gaussian approximation in both accuracy and convergence behavior.

Using Savage-Dickey Density Ratios, we tested and validated the presence of spatial dependence in volatility. The strong Bayes factors in favor of a non-zero spatial autoregressive parameter ( $\lambda$ ) across multiple specifications confirmed that volatility in housing returns is indeed spatially clustered.

Our application to Zillow data found that urban centers generally exhibit higher volatility, while suburban and peripheral areas show more stability—providing empirical evidence for the inefficiencies and heterogeneity in U.S. housing markets. This insight has practical implications for real estate investors, policymakers, and risk managers aiming to understand spatial dynamics in market risk.

In conclusion, our work highlights the importance of modeling higher-order spatial dependencies in volatility, going beyond traditional spatial econometric models that focus primarily on mean dependence. Future extensions may include dynamic spatial volatility models, incorporation of covariates, and multivariate extensions to study cross-market interactions.

## 9 Contribution

- Finding papers: Kunal Dayma, Shivanie
- Understanding methodology : Kunal Dayma, Shivanie
- Coding : Varada Agarwal

- Generating figures and tables : Varada Agarwal
- Writing the report : Kunal Dayma, Shivanie